

Conducting experiments on Mechanical Turk via *mturkutils*

Florian Jaeger with help from Linda Liu, Zach Burchill, Wednesday Bushong, and Xin Xie

October 22, 2021

Contents

1 Overview

Make sure you have completed the tutorial on *Setting up Mechanical Turk and boto(3)*. Now that you are familiar with boto(3) and the mTurkutils, it is time to get started on your own experiment. This tutorial walks you through how to sandbox your experiment on MTurk's sandbox, how to upload the (thoroughly tested) HIT to MTurk, and how to collect your data once all assignments for the HIT have been completed. This tutorial does *not* describe how to program your own experiment.

2 Develop and debug your experiment locally

After you have completed the HTML (and potentially Javascript or other) programming necessary for your experiment, the first recommended step is to debug the experiment locally on your computer. In many cases, you will not even need to install a local webserver on your computer. If your experiment does not record speech, video, or alike but only *presents* audio, video, or other stimuli then you likely will be able to simply open your HTML file on your computer in a browser. This allows you to go through your experiment and debug it. **Only after you have completed local debugging of your experiment, is it time to test whether your experiment also works on MTurk.**

3 Sandbox a HIT

The first step following completed developing and debugging is to *sandbox* your experiment. The MTurk sandbox provides an online testing environment that works like the real MTurk while avoiding that you have to pay for the HIT. You can host a HIT on the MTurk sandbox and then you or other team members can take that HIT in order to debug it. This will also be the first step in your debugging process at which you can test whether the data that you obtain from a completed HIT assignment is in the intended format.

You can walk through the steps described next using your own experiment, or using the short mock experiment contained in this tutorial. This mock experiment is hosted on the HLP Lab web server, and you won't have to modify it at all.

- Using your favorite text editor, check out the contents of example-HIT.yaml in the hits/ folder. This is a YAML configuration file. **YAML configuration files specify the basic information for MTurk HITs** that *mturkutils* will use to upload your HIT(s). This includes the HIT description, the payment, the url that specifies where the experiment is hosted, and more (we'll cover configuration files in a more detail in Section ??).
- Carefully, check the YAML configuration file. This is the moment to check the YAML file you are using for the sandbox, as it likely the blueprint for the YAML file you'll be using once you move from sandbox to production.

Sandbox checklist—*before* uploading HIT(s) to sandbox

- Check whether the payment offered in the experiment matches the experiment offered in the HIT (as specified in the YAML configuration file).
- Check whether you want to batchify your HIT(s). Amazon Mechanical Turk now [charges an additional 20% fee for every HIT with more than 10 assignments](#). You can avoid this fee by splitting your HIT(s) into smaller batches of less than 10 assignments using *mturkutils*, as described in Section ??.
- Make sure that your YAML file(s) specify the right URLs and that those URLs point to HTML files that collect the data you intend to collect:
 - * Check whether your YAML file(s) create HITs for all of the intended conditions of your experiment, and *only* those conditions. Also keep in mind that data for all conditions should be elicited at the same time if you want to be able to state in your write-up that participants were randomly assigned to one of the conditions.
 - * Check whether the YAML file(s) for each HIT point to the correct URLs. If the URL that the YAML file points to does not exist all that you (or a future subject) might see is a blank page.
 - * Make sure that each URL or combination of URL parameters has the intended effect. This might require that you check for each of the HITs whether the stimuli, stimuli order, task, etc. correspond to the condition indicated in the URL / by the URL parameters.

- To upload the HIT to MTurk’s sandbox, run (in the *mturkutils/boto3/* directory):

```
python loadHIT.boto3.py -s -c ../hits/example-HIT.yaml
```

Don’t forget the sandbox tag, otherwise the HIT may end up live on MTurk for other workers to do (and you may end up paying real \$\$\$)! If this was successful, you should see a message similar to the following:

```
[$10,000.00] is the initial balance
You can preview your new HIT at:
https://workersandbox.mturk.com/mturk/preview?groupId=some_long_ID
[$10,000.00] is the final balance
```

You should also notice that a new file (*example-HIT.success.yaml*) was automatically generated in *hits/*. Note that this success file will be overwritten everytime you run *loadHIT* for the same configuration file. While we are in sandbox mode, this does not harm but once the experiment is live this would mean that it will be difficult to manage. **So make sure *not* to delete or overwrite the *example-HIT.success.yaml* file: you will be using it to collect your results.**

By default *mturkutils* will use the AWS credentials specified under your default profile (typically in *./aws/credentials*). **If you would like to upload the HITs under an account that requires different AWS credentials**, you can specify the *-p profile_name* tag, where *profile_name* should refer to a profile in your AWS credentials file. For example:

```
python loadHIT.boto3.py -s -c ../hits/example-HIT.yaml -p lab
```

For this to work, there must be a AWS access key and secret access key below a section called [*profile_name*] in your *./aws/credentials* file. For example, for the above example to work your *./aws/credentials* file should contain the following lines:

```
[lab]
# the HLP lab account
aws_access_key_id = <access key>
aws_secret_access_key = <secret access key>
```

- Go to the URL that was output by loadHIT (https://workersandbox.mturk.com/mturk/preview?groupId=some_long_ID), and **complete the HIT** following the considerations described in the next section.

4 During sandboxing

- To complete your sandbox HIT, you will need to log into your MTurk sandbox worker account. If you see a website telling you that you do not have permission to access the HIT, log into your worker sandbox account, and then visit the URL again.
- While you are sandboxing your experiment, you should note any issues that come up. **If your experiment is using Javascript**, it is highly recommended that you watch the Javascript console right from the beginning of the experiment, and keep a record of any errors or notifications that are shown in the console (e.g., for Chrome, right click anywhere on the webpage and select “Inspect” from the menu; then select the *Console* tab).

5 Obtain and check sandbox results

Congratulations! By this point, you have successfully created a HIT (as a requester) and completed a HIT (as a worker). Now you can obtain and carefully check your results.

- **Download your results from the sandbox.** For example, let’s say we have completed a few assignments the example HIT uploaded in the previous section. You can then pull the results from Mechanical Turk’s sandbox with *mturkutils*:

```
python getResults.boto3.py -s -f ../hits/example-HIT.success.yaml
-r ../hits/example-HIT.results_120120.tab
```

Note that the result file will be overwritten each time you run `getResults`—similar to what `loadHIT` does with the success file. Thus, **it’s recommended to include the date in the filename of the results file**, as in the code example in this section. Finally, if you uploaded the HIT(s) to sandbox under a specific profile, make sure to include that same profile name when you get the results (again by adding `-p profile_name`).

- **Parse your results into an interpretable format.** The output of `getResults` is a tab-separated values (.tab) file, with one row per completed assignment, and one column for each variable that your experiment returned to Mechanical Turk. This is usually not the format we need for data visualization and analysis. For any non-trivial experiment, this .tab file will also be hardly human readable. Since **we should carefully check whether the result file contains all the information that we will need later**, the next step is to parse this file into the format we need. This can be done in Python, R, or whatever scripting language you prefer. In R, for example, `read_tsv()` from the `readr` package (part of the `tidyverse` package) makes reading in tab-separated files easy, and the `dplyr` package (also part of the `tidyverse`) helps with wrangling the data into a suitable format. For example:

```
library(tidyverse)

d = read_tsv("PATH_TO_YOUR_RESULT_FILE/YOUR_RESULT_FILE.tab")
```

- Once you have parse your results into a suitable format, it is time to thoroughly check whether the data contains all the information you will need to visualize and analyze your data.

Sandbox checklist—after obtaining results from sandbox

- Check that the HIT results in data being recorded once the participant (you or your team members) press submit. For example, the experiment might hang.
- If the HIT does record data, make sure that the data is in the intended format and that all variables you will need for your analysis are actually being recorded.
- Finally, check that the *format* of the recorded data is the one you intend. For example:
 - * Are all variables included (e.g., subject and item IDs, trial order information, block information, all dependent variables, all conditions)?
 - * Are the values provided for the variables in the format you intended? Are there unexpected NAs?
 - * *Check your design and counter-balancing.* This usually involves both whether nuisance variables are balanced within conditions and across conditions, whether all combinations of items do actually occur in your data, and whether each combination only occurs as often as intended.

6 Getting more out of your YAML configuration file

In addition to the basic components of the YAML file described above, you can add additional information to the YAML file that *mturkutils* can interpret. This section describes three features that are often helpful.

6.1 One YAML file for many HITs (and lists)

Conveniently *mturkutils* allows you to use a single YAML file to specify many different HITs. Specifically, the URL specified in the YAML file can contain a URL parameter whose value is described by a variable surrounded by curly parentheses. Here is an example of just the URL specification from a YAML file with the URL parameter *list* and the parameter value being specified by the variable *condition*:

```
question:
  url: https://www.hlp.rochester.edu/mturk/YOUR_DIR/exp.html?list={condition}
```

In the same YAML file, we then need to specify values that the variable *condition* can take. This is done in the *input* section of the YAML file following the URL specification. Say for an experiment with a 2-by-2 design crossing condition A and B with short and long exposure:

```
question:
  url: https://www.hlp.rochester.edu/mturk/YOUR_DIR/exp.html?list={condition}
  height: 750
  input:
    - condition: A_short
    - condition: B_short
    - condition: A_long
    - condition: B_long
```

This will create four HITs, one each for the following URLs:

- https://www.hlp.rochester.edu/mturk/YOUR_DIR/exp.html?list=A_short
- https://www.hlp.rochester.edu/mturk/YOUR_DIR/exp.html?list=B_short

- https://www.hlp.rochester.edu/mturk/YOUR_DIR/exp.html?list=A_long
- https://www.hlp.rochester.edu/mturk/YOUR_DIR/exp.html?list=B_long

Of course, this will only have the desired result if the HTML and Javascript code at these URLs knows how to handle the different values for the URL parameter *list*. Note that the number of assignments that you have specified at the top of the YAML file determines how many assignments will be collected from MTurk *for each of the URLs*. **The total number of assignments that this approach results in is thus the # of assignments specified in the YAML file *times* the # of rows specified in the input section of the YAML file** (four in this example).

Unfortunately, it *not* *not* currently possible to specify multiple separate URL parameters through this method. If you're Javascript is using multiple URL parameters, you need to either create multiple separate YAML files for each complete URL (including the combination of URL parameters) or you need to modify your Javascript program to parse the single URL parameter that can be specified using the method described here into the multiple separate parameters expected by the remainder of your Javascript.

6.2 Avoiding the MTurk fee for 10 or more assignments

The same approach described in the previous section can also be used to avoid the 20% fee that Amazon charges on HITs with 10 or more assignments. Simply a) set the number of assignments specified in the YAML file to 9 or less, and b) use the input specification described in the previous section to make as many separate 'lists' (all of which are copies of the same list) as needed to reach the targeted number of participants.

For example, let's say you would like to collect data from 24 participants in each of two conditions A and B. We can then get 8 assignments three times for each of the two conditions:

```

---
title: Speech Perception Experiment. MUST BE NATIVE ENGLISH SPEAKER AND WEAR HEADPHONES.
description: Listen to sentences. Takes approximately 30 minutes.
keywords: psychology,experiment,speech,words
reward: 3.00
assignments: 8
#####
## HIT Timing Properties
#####

# this Assignment Duration value is 60 * 90 = 1.5 hour
assignmentduration: 7200

# this HIT Lifetime value is 60*60*24*5 = 5 days
hitlifetime: 172800

# this Auto Approval period is 60*60*24*15 = 15 days
autoapprovaldelay: 1296000

qualifications:
  builtin:
    # this is a built-in qualification -- user must have > 95\% approval rate
    - qualification: PercentAssignmentsApprovedRequirement
      comparator: GreaterThan
      value: 95
      private: true
    # this is a built-in qualification -- user must be in the United States
    - qualification: LocaleRequirement
      comparator: EqualTo
      locale: US
      private: false

question:
  url: https://www.hlp.rochester.edu/mturk/YOUR_DIR/exp.html?list={condition}
  height: 750
  input:
    # 24 subjects per between-subject condition; 48 subjects total
    - condition: A_1
    - condition: A_2
    - condition: A_3
    - condition: B_1
    - condition: B_2
    - condition: B_3

```

Like described in the previous section, this approach assumes that your the HTML and Javascript that the HIT's url points to will recognize the parameters values for list (A_1, A_2, ...) and select the intended list based on those values.

6.3 Preventing unwanted subjects from taking your experiment

Qualifications can be used to keep people from taking your experiment. For example, we standardly include the following qualifications in the YAML file for our HITs, limiting our experiments to workers who are in the US (as far as MTurk can tell based on their IPs), who have 95% approval ratings from previous HITs they have taken, and who have completed at least 100 HITs:

```

qualifications:
  builtin:
    # this is a built-in qualification -- user must have > 95% approval rate
    - qualification: PercentAssignmentsApprovedRequirement
      comparator: GreaterThan
      value: 95
      private: true
    # this is a built-in qualification -- user must have successfully completed 100 HITs
    - qualification: NumberHITsApproved,
      comparator: GreaterThan,
      value: 100
      private: true
    # this is a built-in qualification -- user must be in the United States
    - qualification: LocaleRequirement
      comparator: EqualTo
      locale: US
      private: true

```

More details on [what built-in qualification types are available and how to use them](#) can be found on Amazon’s website, which also contains some helpful [walk-through examples with code](#). For example, if you would like to specifically recruit workers from all states except Hawaii and Alaska—e.g., because you’d like to keep the time zones restricted—the following combination of qualifications should do the job:

```

qualifications:
  builtin:
    # this is a built-in qualification -- user must be in the United States,
    # but not Hawaii or Alaska
    - qualification: LocaleRequirement
      comparator: EqualTo
      locale: US
      private: true
    - qualification: LocaleRequirement
      comparator: NotEqualTo
      locale: (US, HI)
      private: true
    - qualification: LocaleRequirement
      comparator: NotEqualTo
      locale: (US, AK)
      private: true

```

6.3.1 Premium qualifications

For an additional price tag per assignment, Amazon offers a [wide range of additional qualifications](#), based on workers’ self-reported information. This includes demographic information (age, gender, etc.), political orientation, socio-economic status, daily internet usage, and many many more features.

6.3.2 Custom qualifications

Finally, you can use qualifications to prevent anyone who has previously taken an experiment in the same series from taking another experiment of that type. **This Note that qualifications made on your production MTurk account are not visible during sandboxing (and vice versa). This will cause error messages if you try to include private/custom qualifications in the YAML file for sandboxing.** The upshot of this is that you will typically have two separate YAML configuration files—one

for sandboxing and one for production. Only the latter will contain the additional qualifications described next.

1. [Make an arbitrary qualification on the MTurk website.](#)
2. Assign that qualification to all subjects who have taken the previous experiments you are worried about. This can be done with `assignQual.py`. For details, check out:

```
python assignQual.py -h
```

3. In the YAML file for your new experiment add that workers should *not* have that qualification. I.e., your YAML file should contains the following lines under the qualifications section:

```
custom:
  # Has not done one of my other experiments
  - qualification: ID-OF-QUALIFICATION-YOU-MADE
    comparator: DoesNotExist
    private: true
```

4. Since these qualifications are not visible from the sandbox, you will typically add them to only the YAML file once you move to production. **If you intend to block subjects that have taken previous experiments, remember to include these additional qualifications in the YAML configuration file for production.**

If you're not sure whether you have already assigned the relevant qualifications to all workers, you can [download a csv file with all workers you've previously run along with all the qualifications you have previously assigned to them](#) from your MTurk requester account.

7 Going live!

Before you upload your HIT to MTurk, make sure to go through a final check:

Final checklist

- Have you gone through both the both of the Sandbox checklists?
- Does your YAML configuration file for the live production version match the YAML production file for sandboxing (which you have carefully debugged), *except* for:
 - The number of assignments per HIT (which might have been set to 1 for sandboxing but now should yield the desired number of assignments).
 - **Any additional private/custom qualifications**, such as required to block participants who have taken similar previous experiments from your lab.
- Ethics:
 - Does this experiment fall under a protocol approved by the human subject review board?
 - Is the protocol still valid, i.e., not expired?
 - Have you been added as an experimenter to this protocol?
 - Do you have valid CITI? (certificate that you have gone through RSRB training)
 - Does your experiment provide a contact (email) address?
- In the Javascript (.js) file for the experiment:
 - Is the consent form up to date?
 - Are you using the right RSRB protocol number?
 - Is the payment (*Reward*) specified in the YAML file correct?
 - Does the payment (*Reward*) specified in the YAML file match the payment specified in the instructions of your HTML file?
 - Has any code been removed that was meant for testing? (best not to ever add such code)
- **Are there sufficient funds in your MTurk account?** You can check your account balance at <https://requester.mturk.com/account>. If you have insufficient funds, you will receive an error message for all or some of the HITs you are aiming to upload. You can add funds to the account your are using at <https://requester.mturk.com/prepayments/new>.

Once you have completed this checklist, you are ready to upload your HIT to MTurk. Simply repeat the steps outlined above for sandboxing a HIT but without the sandbox tag (-s or -sandbox).

7.1 While the experiment is live

Monitoring while the experiment is in progress

- Check progress of your experiment. If nobody is completing assignments for the HIT, it is likely that something went wrong.
- Regularly check the email account provided as contact for the experiment for emails from MTurk workers. Workers might run into technical difficulties. If that happens, consider the following questions. The sooner you act, the better your chances to create a bad reputation for the lab, to cause frustration among workers, or to waste money and time:
 - Is the problem so severe and likely to apply to others that you should abort the experiment?
 - What can be learned from this problem for future experiments? Keep notes in your to-do file that avoid the same problem in future experiments.
 - Be responsive to workers concerns and complaints. See Section ?? for a list of common issues that come up during experiments.
 - Offer to pay people that have tried to complete a HIT but then ran into technical difficulties (e.g., the computer hung). See Section ??

7.2 Common issues, how to avoid, and how to address them

A few issues tend to come up in almost every experiment, most of them relating to technical difficulties or workers not reading your instructions in the way you intended.

- Workers are sometimes concerned that they won't get paid because they misclicked or mistyped on a few trials. Consider rewording your instructions so that it is clear that occasional mistakes are not only understandable (and will still result in payment) but are, in fact, of interest to researchers in the cognitive sciences. Be clear that whatever performance measures you are eliciting are there to exclude the *very few* workers who actually game the system with random clicking.
- A common reason for technical difficulties are browser incompatibilities. What you can do is a) try to use javascript that is robust across browsers and operating systems. Generally basic javascript is better than using packages, though even basic javascript is sometimes interpreted differently across environments; b) test / sandbox your experiment across multiple platforms; c) recognize incompatible browser types in your javascript, and prevent those user from taking your experiment; d) **state clearly in the description of your HIT which browsers and operating systems are acceptable.**
- Workers tend to contact you stating that something didn't work—the experiment froze, they had to reload it, they couldn't submit the HIT. Our first question is usually, please take a screen shot, tell us your browser version and operating system, and (for technically advanced workers) send us a copy of the console output. However, often that information is not available by the time workers contact you, as they have typically closed that browser window. **It's a good idea to alert workers in your instructions to take screen shots of problems.**
- If you use cookies to prevent workers from taking the experiment twice, that can create confusion, also because some people share the same worker accounts. Our answer to those cases is always that we cannot ascertain that it is not the same individual, and the integrity of our research would not be guaranteed if we allowed the same individual to take the same experiment more than once. I also recommend writing your instructions in such a way that this is very clear.

7.3 Getting your results

Follow the steps described above for getting results from the sandbox except without the `-s` flag.

8 Soon after the experiment

8.1 Paying (approving) workers

If you don't pay workers quickly, then you start getting bad feedback and no one will want to do your HITs. The YAML configuration file for your experiment specifies a period after which workers who have completed an assignment for the HIT(s) will be automatically approved. It's this part in the YAML file:

```
# this Auto Approval period is 60*60*24*15 = 15 days
autoapprovaldelay: 1296000
```

However, it's best to think of the autoapprove as an insurance policy in case you forget. You want the auto-approve interval long enough time to let yourself reject bad actors' work (if you can) but you want to approve good work quickly enough that you get good ratings on the turker boards. To pay all the workers contained in a result file, use:

```
python approveWork.boto3.py -r RESULT_FILE -p PROFILE_NAME
```

IRB rules dictate that anyone taking the experiment should be reimbursed. This means you should approve *all* workers who have taken your experiment.

8.2 Blocking workers from *any* future experiments

After reimbursement though you might be interested in excluding participants who *clearly* gamed the system from future experiments (see next section). Be careful though: ask yourself whether it wasn't lack of clarity of the instructions, rather than nefarious intent on the participant's side that caused the pattern of data. In our experience, the former is the far more common cause for data that appears to suggest that participants are not really doing the task you intended. To quote from Amazon's "Best practices" (emphasis added):

A word of caution: MTurk has a built-in capability for a Requester to "block" a Worker from working on that Requester's tasks in the future. Requesters use this capability to identify and exclude Workers that have consistently produced poor quality results on their tasks. **We strongly urge Requesters to avoid using this feature to exclude Workers for convenience. Doing so could negatively impact the Worker's reputation in MTurk.**

If, after careful consideration, you decide that you want to permanently block a worker from *all* experiment run through your requester account, you can use the following *mturkutils* script:

```
python blockWorkers.boto3.py -h
```

8.3 Blocking workers from *similar* future experiments

Additionally, you can exclude workers from *specific* future experiments. As discussed in Section ??, you can create a new private/custom qualification on your MTurk requester account that identifies all participants who have taken your experiment. For that you will need a result file with worker IDs. All of the IDs in that file will be assigned the new qualification. For more details, see:

```
python assignQual.boto3.py -h
```

8.4 Paying workers who couldn't complete the experiment due to technical difficulties

IRB rules dictate that any participant has the right to abort an experiment at any point. We interpret this to mean that one should approve every worker who has taken the experiment, including those who ran into technical difficulty or aborted prior to completion of the assignment. **The easiest way to pay these workers** is:

- Set up a new qualification through the MTurk requester website (see Section ??).
- Assign that qualification to the workers in question. (If you are not sure whether you have already done so, you can [download a csv file with all workers you've previously run along with all the qualifications you have previously assigned to them](#) from your MTurk requester account.)
- Make a new HIT that just contains a single question like “Click this checkbox to get paid and press submit.”
- Send out an email to the worker(s) and tell them to search for that HIT (since one apparently can't directly send links to HITs).

The following is a screen shot of a HIT description that was used for remuneration.

[Manage Batches](#) > [Batch Details](#)

Chinese Sentence Renumeration HIT 1

View the latest status of this batch, make changes, or get results.

Get compensated for your time despite technical difficulties

Status

Status: Pending Review

100% submitted

100% published

Assignments Completed: 1 / 1

Creation Time: January 22, 2017 7:24 AM PST

Average Time per Assignment: 13 seconds

Completion Time: January 22, 2017 3:02 PM PST

Settings

Chinese Sentence Renumeration HIT

[View Project](#)

Note: If you have edited the Project after publishing this Batch, you will see the latest version.

Description: Get compensated for your time despite technical difficulties

Keywords:

Qualification Requirement(s): Chinese Sentence Reading Remuneration has been granted

Number of Assignments per task: 1

Reward per Assignment: \$1.50

Batch expired on: January 29, 2017 7:24 AM PST

Assignment duration: 20 minutes

Auto Approval Delay: 8 hours

Results

Assignments pending review: 0

Assignments approved: 1

Assignments rejected: 0

Cost Summary

Estimated Total Reward: \$1.50

Estimated Fees to Mechanical Turk: \$0.30 ([fee details](#))

Estimated Total Cost: \$1.80

These costs are only an estimate until all of the assignments have been submitted and reviewed.

Example task from this Batch

Write anything below

You must ACCEPT the HIT before you can submit the results.

Figure 1: Screenshot of a HIT made specifically to remunerate workers who could not complete the experiment because of technical difficulties.

8.5 Granting bonuses

In addition to payments, MTurk provides the option to grant a bonus to a worker. This option can come in handily, for example, when there are multiple visits to an experiment (e.g., a longitudinal study). In that case the reward would be the basic payment that the worker received for completing any part of the experiment, and the bonus would be used as an incentive to complete *all* visits of the experiment. *Mturkutils* comes with a separate script to grant bonuses to workers.